

YIFAN JIANG

yjiang44@usc.edu/yifjia@isi.edu
https://yifanjiang-921.github.io/

Education Background

University of Southern California (USC)

Ph.D. in Computer Science

- Advisor: Dr. Jay Pujara

Los Angeles, USA

08/2023~Current

University of Southern California (USC)

M.S. in Computer Science

Overall GPA: 3.9/4.0

- Advisor: Dr. Filip Ilievski

Los Angeles, USA

01/2022~12/2023

Southern University of Science and Technology (SUSTech)

B.S. in Computer Science

Overall GPA: 3.78/4.0

- Advisor: Dr. Xuan Song

Shenzhen, China

09/2017~06/2021

University of California, Berkeley (UCB)

Exchange program in Computer Science

Overall GPA: 4.0/4.0

San Francisco, USA

01/2020~05/2020

Experience

Amazon Prime Video

Research Intern

Seattle, USA

05/2025~Current

Hippocratic AI

Research Intern

Palo Alto, USA

05/2024~08/2024

Information Sciences Institute, USC

Research Assistant

Los Angeles, USA

05/2022~08/2023

Computer Science Department, USC

Teaching Assistant

Los Angeles, USA

06/2022~12/2022

SUSTech-Tokyo Joint Research Center for Super Smart City

Research Assistant

Shenzhen, China

01/2021~12/2021

Publication

- Engage for All: Making Ordinary Image Descriptions Appealing Again!
Yuyan Chen, **Yifan Jiang**, Li Zhou, Jinghan Cao, Yu Guan, Ming Yang, Qingpei Guo
[ICCV 2025] International Conference on Computer Vision
- RED QUEEN: Safeguarding Large Language Models against Concealed Multi-Turn Jailbreaking
Yifan Jiang, Kriti Aggarwal, Subhabrata Mukherjee, Tanmay Laud, Kashif Munir, Jay Pujara
[ACL 2025 Findings] The 63rd Annual Meeting of the Association for Computational Linguistics
- COLUMBUS: A Visual Question Answering Benchmark to Evaluate COgnitive Lateral Understanding through Multiple-choice reBUSes
Koen Kraaijveld, **Yifan Jiang**, Kaixin Ma, Filip Ilievski
[AAAI 2025] AAAI Conference on Artificial Intelligence, 2025
- MARVEL: Multidimensional Abstraction and Reasoning through Visual Evaluation and Learning
Yifan Jiang^{*}, Jiarui Zhang^{*}, Kexuan Sun^{*}, Zhivar Sourati, Kian Ahrabian, Kaixin Ma, Filip Ilievski, Jay Pujara
[NeurIPS DB 2024] Conference on Neural Information Processing Systems, 2024

- The curious case of nonverbal abstract reasoning with multi-modal large language models
Kian Ahrabian*, Zhivar Sourati*, Kexuan Sun*, Jiarui Zhang, **Yifan Jiang**, Fred Morstatter, Jay Pujara
[COLM 2024] Conference On Language Modeling, 2024
- ARN: Analogical Reasoning on Narratives
Zhivar Sourati, Filip Ilievski, Pia Sommerauer, **Yifan Jiang**
[TACL 2024] Transactions of the Association for Computational Linguistics, 2024.
- Semeval-2024 task 9: Brainteaser: A novel task defying common sense
Yifan Jiang, Filip Ilievski, Kaixin Ma
[SemEval@NAACL 2024] The 18th International Workshop on Semantic Evaluation, 2024.
- BRAINTEASER: Lateral Thinking Puzzles for Large Language Models
Yifan Jiang, Filip Ilievski, Kaixin Ma, Zhivar Sourati
[EMNLP 2023] Conference on Empirical Methods in Natural Language Processing, 2023.
- Transferring Procedural Knowledge across Commonsense Tasks
Yifan Jiang, Kaixin Ma, Filip Ilievski
[WNU@ACL 2023] ACL 2023 5th Workshop on Narrative Understanding.
[ECAI 2023] European Conference on Artificial Intelligence, 2023.
- FIRE: Food Image to REcipe generation
Prateek Chhikara, Dhiraj Chaurasia, **Yifan Jiang**, Omkar Masur
[WACV 2024] IEEE/CVF Winter Conference on Applications of Computer Vision, 2024.
- TrafPS: A Shapley-based Visual Analytics Approach to Interpret Traffic
Zezheng Feng, **Yifan Jiang**, Hongjun Wang, Zipei Fan, Yuxin Ma, Shuang-Hua Yang, Huamin Qu, Xuan Song
[CVMJ 2023] Computational Visual Media Journal, 2023.
- SRPO: Enhancing multimodal llm reasoning via reflection-aware reinforcement learning
Zhongwei Wan, Zhihao Dou, Che Liu, Yu Zhang, Dongfei Cui, Qinjian Zhao, Hui Shen, Jing Xiong, Yi Xin,
Yifan Jiang, Yangfan He, Mi Zhang, Shen Yan
[Preprint]
- Recent Advances in Large Language Model Benchmarks against Data Contamination: From Static to Dynamic Evaluation
Simin Chen, Yiming Chen, Zexin Li, **Yifan Jiang**, Zhongwei Wan, Yixin He, Dezhi Ran, Tianle Gu, Haizhou Li, Tao Xie, Baishakhi Ray
[Preprint]
- A Comprehensive Survey in LLM (-Agent) Full Stack Safety: Data, Training and Deployment
[Preprint]

Service

- [2025] Reviewer for COLIING, COLING_demo, ICLR, IEEE_SSCI, Analogy-ANGLE Work shop II
- [2024] Reviewer for SemEval, LREC-COLING, Analogy-ANGLE Workshop(co-located with IJCAI)
- [2025] Reviewer for EMNLP, NeurIPS.

Honors

- The Honor Roll Merit Scholarship, Southern University of Science and Technology, 2018,2019,2020
- Computer Science Outstanding Graduate Award, Southern University of Science and Technology, 2021
- Computer Science Honor Program, University of Southern California, 2022
- ISI Distinguished Graduate Student Researcher Fellowship Award, University of Southern California, 2023